





GAN

Name	Nathan Varghese
Identity Key	nawa3000

	Level	Completed	Goal	
	Beginner	4	5722	15
	Intermediate	3	Total Completed	
	Advanced	2		
	Expert	0		
			9	

Generative Adversarial Network (GAN)

CSCI 5277: Computer Vision

Fall 2024

Dr. Tom Yeh



2014



2015



2016



2017



2018



2019



2020



2021



2022



2023

Generative Adversarial Nets

Ian J. Goodfellow*, Jean Pouget-Abadie†, Mehdi Mirza, Bing Xu, David Warde-Farley,
Sherjil Ozair‡, Aaron Courville, Yoshua Bengio§

Département d'informatique et de recherche opérationnelle
Université de Montréal
Montréal, QC H3C 3J7

Abstract

We propose a new framework for estimating generative models via an adversarial process, in which we simultaneously train two models: a generative model G that captures the data distribution, and a discriminative model D that estimates the probability that a sample came from the training data rather than G . The training procedure for G is to maximize the probability of D making a mistake. This framework corresponds to a minimax two-player game. In the space of arbitrary functions G and D , a unique solution exists, with G recovering the training data distribution and D equal to $\frac{1}{2}$ everywhere. In the case where G and D are defined by multilayer perceptrons, the entire system can be trained with backpropagation. There is no need for any Markov chains or unrolled approximate inference networks during either training or generation of samples. Experiments demonstrate the potential of the framework through qualitative and quantitative evaluation of the generated samples.

1 Introduction

The promise of deep learning is to discover rich, hierarchical models [2] that represent probability distributions over the kinds of data encountered in artificial intelligence applications, such as natural images, audio waveforms containing speech, and symbols in natural language corpora. So far, the most striking successes in deep learning have involved discriminative models, usually those that map a high-dimensional, rich sensory input to a class label [14, 20]. These striking successes have primarily been based on the backpropagation and dropout algorithms, using piecewise linear units [17, 8, 9] which have a particularly well-behaved gradient. Deep *generative* models have had less of an impact, due to the difficulty of approximating many intractable probabilistic computations that arise in maximum likelihood estimation and related strategies, and due to difficulty of leveraging the benefits of piecewise linear units in the generative context. We propose a new generative model estimation procedure that sidesteps these difficulties.¹

In the proposed *adversarial nets* framework, the generative model is pitted against an adversary: a discriminative model that learns to determine whether a sample is from the model distribution or the data distribution. The generative model can be thought of as analogous to a team of counterfeiters, trying to produce fake currency and use it without detection, while the discriminative model is analogous to the police, trying to detect the counterfeit currency. Competition in this game drives both teams to improve their methods until the counterfeits are indistinguishable from the genuine articles.

*Ian Goodfellow is now a research scientist at Google, but did this work earlier as a UdeM student

†Jean Pouget-Abadie did this work while visiting Université de Montréal from Ecole Polytechnique.

‡Sherjil Ozair is visiting Université de Montréal from Indian Institute of Technology Delhi

§Yoshua Bengio is a CIFAR Senior Fellow.

¹All code and hyperparameters available at <http://www.github.com/goodfeli/adversarial>



Share your AI-generated image of an ANN

<https://www.linkedin.com/feed/update/urn:li:activity:7241258998543314945>



Computer Vision - S24

Tom Yeh • You
2m • Edited •

Activity: Share your AI-generated image of an ANN

Post a comment with your image, the prompt you provided, and the tool you

1 comment



Like



Comment



Add a comment...



Most recent ▾



Tom Yeh (He/Him) **Author**

2m (edited) ...


Associate Professor of Computer Science at University of Colorado B...

Prompt:

Generate a graphic representation of an artificial neural network

Tool:

ChatGPT4



Like | Reply

Visit the link above
Add a comment
Attach your image



Nathan Varghese

MS Aerospace Engineering at CU Boulder | Autonomous Systems

Prompt: Draw a graphic representation of a artificial neural network. I want you to show it as a complicated spaghetti-fied image.



Like | Reply

GAN Architecture

CSCI 5722 Computer Vision



University of Colorado
Boulder

Architecture Diagram





Diagram → ANN

Random
○○○




Fake
○○○



Real?

Real
○○○



Random

1	1	-1
1	-1	1

1 1 1

Fake

-1	-1	-3
3	-1	3

1 1 1

Real

0	0	-1
3	-1	1

1 1 1

1 1 1

1 1 1

--	--	--

--	--	--

Pred

.3	.2	.4
----	----	----

.5	.3	.7
----	----	----

Generator


CSCI 5722 Computer Vision



University of Colorado
Boulder

Generator

			A	B	C	
			1	1	-1	
			1	-1	1	
			1	1	1	
1	0	-2	-1	-1	-3	x
0	2	1	3	-1	3	y



Discriminator


CSCI 5722 Computer Vision



University of Colorado
Boulder

Discriminator

			A	B	C	D	E	F
			-1	-1	-3	0	0	-1
			3	-1	3	3	-1	1
			1	1	1	1	1	1
W_1	1	0	1					
	0	1	1					
			1	1	1	1	1	1
W_2	1	1	-3					





$$x \geq 1$$

$$x - 1 \geq 0$$

$$x \geq -1$$

$$x + 1 \geq 0$$



	x	y
V_1	1	0
V_2	1	0

Fake Real

A	B	C	D	E	F
-1	-1	-3	0	0	-1
3	-1	3	3	-1	1
1	1	1	1	1	1



ReLU

$$y \geq 1$$

$$y-1 \geq 1$$

$$y \geq 1$$

$$y+1 \geq 0$$



	A	B	C	D	E	F
	-1	-1	-3	0	0	-1
	3	-1	3	3	-1	1
	1	1	1	1	1	1
h_1	0	1	-1	2	2	0
h_2	0	1	1	4	0	2

ReLU




Activation Field (ReLU)

V: $x = -1$

H: $y = -1$



$$a_1 = \text{ReLU}(x+1)$$

$$a_2 = \text{ReLU}(y+1)$$



☒ ☐ Activation Field ($b = 0$)

L: $x+y = 0$ a: $\text{ReLU}(x+y)$



sum(a) = 7; sum(b) = 11; sum(c) = 3



(Drag L to the right)



Activation Field ($b = -3$)

$$L: x+y = 3$$

$$a: \text{ReLU}(x+y-3)$$



$$x+y-3=0$$

$$2+y-3=0$$

$$y=+1$$



$\text{sum}(a) = 8 ; \text{sum}(b) = 0$



Weighted Sum Field

$$L: 2x + y = 0$$

a	-3	-2	-1	0	1	2	3	
3	-3	-1	1	3	5	7	9	c
2	-4	-2	0	2	4	6	8	
1	-5	-3	-1	1	3	5	7	
0	-6	-4	-2	0	2	4	6	b
-1	-7	-5				3	5	
-2			-4	-2	0			
-3			-5	-3	-1			

$$\begin{array}{r} 6 + 7 = 13 \\ 8 + 9 = 17 \\ \hline 0 \end{array}$$

L	2	1	0	?
---	---	---	---	---

*
*


1

$$2 \times 2 + 3 = 7$$

$$\begin{aligned} 2(-1) + y &= 0 \\ -2 &= -y \\ y &= 2 \end{aligned}$$

$$2(-2)$$

$$2x + 1$$




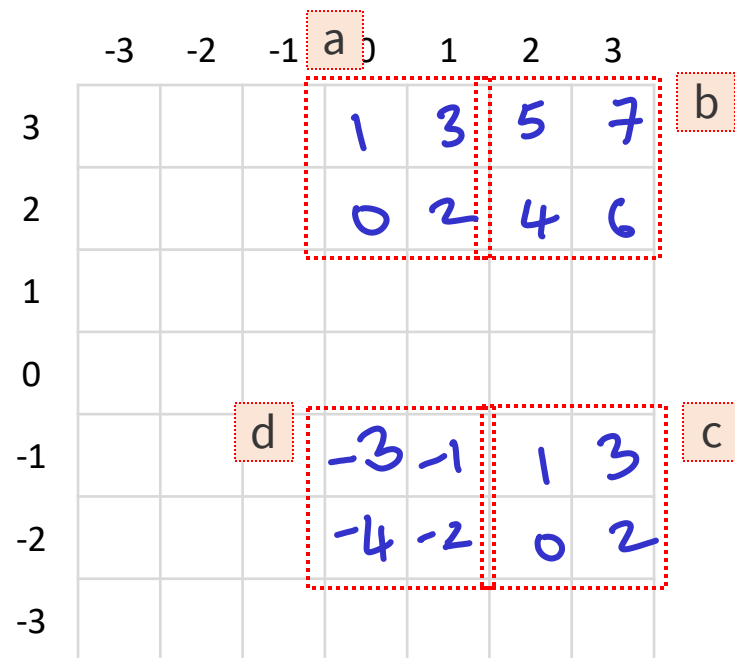
$$\begin{aligned} 2x + 1 &= 0 \\ x &= -1/2 \end{aligned}$$



sum(a) = -10; sum(b) = 3; sum(c) = 30;

Activation Field ($b = -2$)

$L: 2 * x + y = 2$ $a: \text{ReLU}(2 * x + y - 2)$



$$2x + y - 2 = 0$$

$$2(0) + y - 2 = 0$$

$$y = 2$$

$$0 \times 2 + 3 - 2 = 1$$

$$2 + 3 - 2 = 3$$

$$4 + 3 - 2 = 5$$

$$6 + 3 - 2 = 7$$

$$0 \times 2 + 2 - 2 = 0$$


$$2 + 2 - 2 = 2$$

$$0 - 1 - 2 = -3$$



$$0 - 2 - 2 = -4$$

$$2 - 1 - 2 = -1$$

$$2 - 2 - 2 = -2$$

 $\text{sum}(a) = 6; \text{sum}(b) = 22; \text{sum}(c) = 6;$

Logit Field ($b = 0$)

$$V_2$$
$$\text{H:}_2$$



Logit Field (b = -3)

V: $x = -1$


a_1 : $\text{ReLU}(x+1)$

H: $y = -1$

a_2 : $\text{ReLU}(y+1)$




$$z = a_1 + a_2$$




☒ ☐ Logit Field (b = 0)

V: x = 1

H: y = 1



(Drag H, V to the right places)



$$\begin{array}{r}
 2-1=1 \\
 3-1=2 \approx 3 \\
 \hline
 3-1=2 \\
 3-1=2 \approx 4 \\
 \hline
 2-1=1 \\
 2-1=1 = 2 \\
 \hline
 3-1=2 \\
 2-1=1 = 3 \\
 \hline
 2-1=1 \\
 1-1=0 \} = 1 \\
 \hline
 1-1=0 \\
 0-1=-1 \} = -1 \\
 \hline
 -1
 \end{array}$$

sum(a) = 12; sum(b) = 6;



Logit Field


$$3 - 2 = 1$$

$$= 5$$

$$L_1: x+y = 2 \quad a_1: \text{ReLU}(x+y-2) \quad z: a_1+a_2$$

$$L_2: x+y = -2 \quad a_2: \text{ReLU}(-x-y-2)$$



 $\text{sum}(a) = 4; \text{sum}(b) = 12; \text{sum}(c) = 4$



Logit Field

L: $x+y = -1$

a_1 : $\text{ReLU}(x+y+1)$



z : a_1+a_2

V: $x = 1$

a_2 : $\text{ReLU}(x-1)$

Handwritten calculations:

$$\begin{array}{r} 0+3+1=4 \\ -1 \sim -1 \\ \hline 1+3+1=5 \\ 1-1=0 \\ \hline 0+2+1=3 \\ =-1 \end{array} \quad \left. \vphantom{\begin{array}{r} 0+3+1=4 \\ -1 \sim -1 \\ \hline 1+3+1=5 \\ 1-1=0 \\ \hline 0+2+1=3 \\ =-1 \end{array}} \right\} a$$



key icon $\text{sum}(a) = 16$; $\text{sum}(b) = 30$

Logit Field \rightarrow Output Probability Field

V: $x = -1$


a_1 : $\text{ReLU}(x+1)$

$z = a_1 + a_2 - 3$


H: $y = -1$

a_2 : $\text{ReLU}(y+1)$

		-3	-2	-1	0	1	2	3
3		1	1	1	2	3	4	5
2		0	0	0	1	2	3	4
1		-1	-1	-1	0	1	2	3
0		-2	-2	-2	-1	0	1	2
H	-1	-3	-3	-3	-2	-1	0	1
-2		-3	-3	-3	-2	-1	0	1
-3		-3	-3	-3	-2	-1	0	1
				V				



Approximate Sigmoid Function by Hand 🖍️



$$\sigma(\leq -3) \approx$$

$$\sigma(-2) \approx$$

$$\sigma(-1) \approx$$

$$\sigma(0) \approx$$

$$\sigma(-1) \approx$$


$$\sigma(-2) \approx$$

$$\sigma(\geq 3) \approx$$

Output Probability Field

	-3	-2	-1	0	1	2	3
3	1	1	1	2	3	4	5
2	0	0	0	1	2	3	4
1	-1	-1	-1	0	1	2	3
0	-2	-2	-2	-1	0	1	2
-1	-3	-3	-3	-2	-1	0	1
-2	-3	-3	-3	-2	-1	0	1
-3	-3	-3	-3	-2	-1	0	1

$\sigma(\leq -3) \approx 0$
 $\sigma(-2) \approx 0.1$
 $\sigma(-1) \approx 0.3$
 $\sigma(0) \approx 0.5$
 $\sigma(-1) \approx 0.7$
 $\sigma(-2) \approx 0.9$
 $\sigma(\geq 3) \approx 1$



GAN

Noise		
1	1	-1
1	-1	1
1	1	1
Fake		
A	B	C
1	0	-2
0	2	1
1	1	1
Real		
D	E	F
-1	-1	-3
3	-1	3
1	1	1
[ReLU]		
1	1	-3
0.7	0	0.7
6.7	0.1	6.3

df
prob

C		A	D			
		F				
		B	E			

.7	.7	.7	.9	1	1	1
.5	.5	.5	.7	.9	1	1
.3	.3	.3	.5	.7	.9	1
.1	.1	.1	.3	.5	.7	.9
0	0	0	.1	.3	.5	.7
0	0	0	.1	.3	.5	.7
0	0	0	.1	.3	.5	.7

Binary Cross Entropy (BCE) Loss

CSCI 5722 Computer Vision



University of Colorado
Boulder

Categorical \rightarrow Binary CE

!dog

(


	P ₁	P ₂	P ₃	P ₄	P ₅
dog					
cat					
cow					
pig					

Σ

T

!dog
dog cat cow pig

P ₁				
P ₂				
P ₃				
P ₄				
P ₅				



Q₁ Q₂ Q₃ Q₄ Q₅

dog	0	0	.9	.1	1
cat	0	.9	.1	.9	0
cow	.1	.1	0	0	0
pig	.9	0	0	0	0

Σ


-log Q

$\Sigma p - \log Q$

$-\log(0) \approx 32$
 $-\log(-1) \approx 3$
 $-\log(0.9) \approx 0.2$
 $-\log(1) = 0$

Binary CE Loss

	P ₁	P ₂	P ₃	P ₄	P ₅		Q ₁	Q ₂	Q ₃	Q ₄	Q ₅
dog	0	0	0	0	1	dog	0	0	.9	.1	1
!dog	1	1	1	1	0	!dog					



Q₁ Q₂ Q₃ Q₄ Q₅

0	0	.9	.1	1

\hat{y}
 $1-\hat{y}$

$-\log(0) \approx 32$
 $-\log(0.1) \approx 3$
 $-\log(0.9) \approx 0.2$
 $-\log(1) = 0$

T

dog !dog

P ₁	0	1
P ₂	0	1
P ₃	0	1
P ₄	0	1
P ₅	1	0

y $(1-y)$


$-\log(\hat{y})$
 $-\log(1-\hat{y})$

Binary Cross entropy Loss

$$y \log(\hat{y}) + (1-y) \log(1-\hat{y})$$

Math: BCE Loss

	P ₁	P ₂	P ₃	P ₄	P ₅		Q ₁	Q ₂	Q ₃	Q ₄	Q ₅
dog	0	0	0	0	1	dog	0	0	.9	.1	1
!dog	1	1	1	1	0	!dog	1	1	0.1	0.9	0.1



	dog	!dog
P ₁	0	
P ₂	0	
P ₃	0	
P ₄	0	
P ₅	1	

				0
0	0	3	0.2	



BCE Loss

	P ₁	P ₂	P ₃	P ₄	P ₅
pig	0	0	1	1	0
!pig	1	1	0	0	1

	Q ₁	Q ₂	Q ₃	Q ₄	Q ₅
pig	0	1	.9	.1	0
!pig	1	0	0.1	0.9	1

$-\log(0) \approx 32$
 $-\log(0.1) \approx 3$
 $-\log(0.9) \approx 0.2$
 $-\log(1) = 0$

		.	.	0.2	3	.
		0	32	.	.	0
	a	0	b			
P ₁	0	1				
P ₂	0	1	32	c		
P ₃	1	0		0.2	d	
P ₄	1	0			3	e
P ₅	0	1				6

🔑 $a+b+c+d+e =$
35.2;

GAN's Loss

CSCI 5722 Computer Vision



University of Colorado
Boulder

Discriminator Loss

P_1	P_2	P_3	P_4	P_5	P_6

P_1		
P_2		
P_3		
P_4		
P_5		
P_6		

Fake			Real		
-1	-1	-3	0	0	-1
3	-1	3	3	-1	1

.7	0	.7	.9	.1	.3
----	---	----	----	----	----

$-\log(0) \approx 32$
 $-\log(0.1) \approx 3$
 $-\log(0.3) \approx 2$
 $-\log(0.5) \approx 1$
 $-\log(0.7) \approx 0.5$
 $-\log(0.9) \approx 0.2$
 $-\log(1) = 0$

Generator Loss

P ₁	P ₂	P ₃	P ₄	P ₅	P ₆

P ₁		
P ₂		
P ₃		
P ₄		
P ₅		
P ₆		

Fake			Real		
-1	-1	-3	0	0	-1
3	-1	3	3	-1	1


.7	0	.7	.9	.1	.3
----	---	----	----	----	----

$-\log(0) \approx 32$
 $-\log(0.1) \approx 3$
 $-\log(0.3) \approx 2$
 $-\log(0.5) \approx 1$
 $-\log(0.7) \approx 0.5$
 $-\log(0.9) \approx 0.2$
 $-\log(1) = 0$



Discriminator Loss

	P ₁	P ₂	P ₃	P ₄	P ₅	P ₆
real	0	0	0	0	1	1
!real	1	1	1	1	0	0





.3	.7	.5	.3	.1	1

Y
1-Y

-log

	real	!real
P ₁	0	1
P ₂	0	1
P ₃	0	1
P ₄	0	1
P ₅	1	0
P ₆	1	0

	a				
		b			
			c		
				d	
					e
					f



$-\log(0) \approx 32$
 $-\log(0.1) \approx 3$
 $-\log(0.3) \approx 2$
 $-\log(0.5) \approx 1$
 $-\log(0.7) \approx 0.5$
 $-\log(0.9) \approx 0.2$
 $-\log(1) = 0$


$a+b+c+d+e+f = 7;$



Generator Loss

	P ₁	P ₂	P ₃	P ₄	P ₅	P ₆
real	1	1	1	1	0	0
!real	0	0	0	0	0	0

	real	!real
P ₁	1	0
P ₂	1	0
P ₃	1	0
P ₄	1	0
P ₅	0	0
P ₆	0	0





.3	.7	.5	.3	.1	1

Y
1-Y

-log

	a				
		b			
			c		
				d	
					e
					f




$-\log(0) \approx 32$
 $-\log(0.1) \approx 3$
 $-\log(0.3) \approx 2$
 $-\log(0.5) \approx 1$
 $-\log(0.7) \approx 0.5$
 $-\log(0.9) \approx 0.2$
 $-\log(1) = 0$


$a+b+c+d+e+f = 5.5;$




Discriminator Loss



Fake



Real



$-\log(0) \approx 32$
 $-\log(0.1) \approx 3$
 $-\log(0.3) \approx 2$
 $-\log(0.5) \approx 1$
 $-\log(0.7) \approx 0.5$
 $-\log(0.9) \approx 0.2$
 $-\log(1) = 0$

	P ₁	P ₂	P ₃	P ₄	P ₅	P ₆
real						
!real						

.3	.1	.5	.9	.1	0

Y
1-Y

-log

	real	!real
P ₁		
P ₂		
P ₃		
P ₄		
P ₅		
P ₆		

	a				
		b			
			c		
				d	
					e
					f


$a+b+c+d+e+f=$
36.9;



Generator Loss

	P ₁	P ₂	P ₃	P ₄	P ₅	P ₆
real	1		1		1	
!real						

	real	!real
P ₁	1	
P ₂		
P ₃	1	
P ₄		
P ₅	1	
P ₆		





.3	.1	.5	.9	.1	0

Y
1-Y

-log

	a				
		b			
			c		
				d	
					e
					f



- log(0) ≈ 32
- log(0.1) ≈ 3
- log(0.3) ≈ 2
- log(0.5) ≈ 1
- log(0.7) ≈ 0.5
- log(0.9) ≈ 0.2
- log(1) = 0

🔑 $a+b+c+d+e+f = 6;$


BCE Loss Gradients

CSCI 5722 Computer Vision



University of Colorado
Boulder

Update a parameter




1	1	1	2	3	4	5
0	0	0	1	2	3	4
-1	-1	-1	0	1	2	3
-2	-2	-2	-1	0	1	2
-3	-3	-3	-2	-1	0	1
-3	-3	-3	-2	-1	0	1
-3	-3	-3	-2	-1	0	1

.7	.7	.7	.9	1	1	1
.5	.5	.5	.7	.9	1	1
.3	.3	.3	.5	.7	.9	1
.1	.1	.1	.3	.5	.7	.9
0	0	0	.1	.3	.5	.7
0	0	0	.1	.3	.5	.7
0	0	0	.1	.3	.5	.7

Fake			Real		
.7	0	.7	.9	.1	.3

Discriminator Loss Gradients



-log


Target

--	--	--	--	--	--

$\sigma(\leq -3) \approx 0$
 $\sigma(-2) \approx 0.1$
 $\sigma(-1) \approx 0.3$
 $\sigma(0) \approx 0.5$
 $\sigma(1) \approx 0.7$
 $\sigma(2) \approx 0.9$
 $\sigma(\geq 3) \approx 1$

$-\log(0) \approx 32$
 $-\log(0.1) \approx 3$
 $-\log(0.3) \approx 2$
 $-\log(0.5) \approx 1$
 $-\log(0.7) \approx 0.5$
 $-\log(0.9) \approx 0.2$
 $-\log(1) = 0$

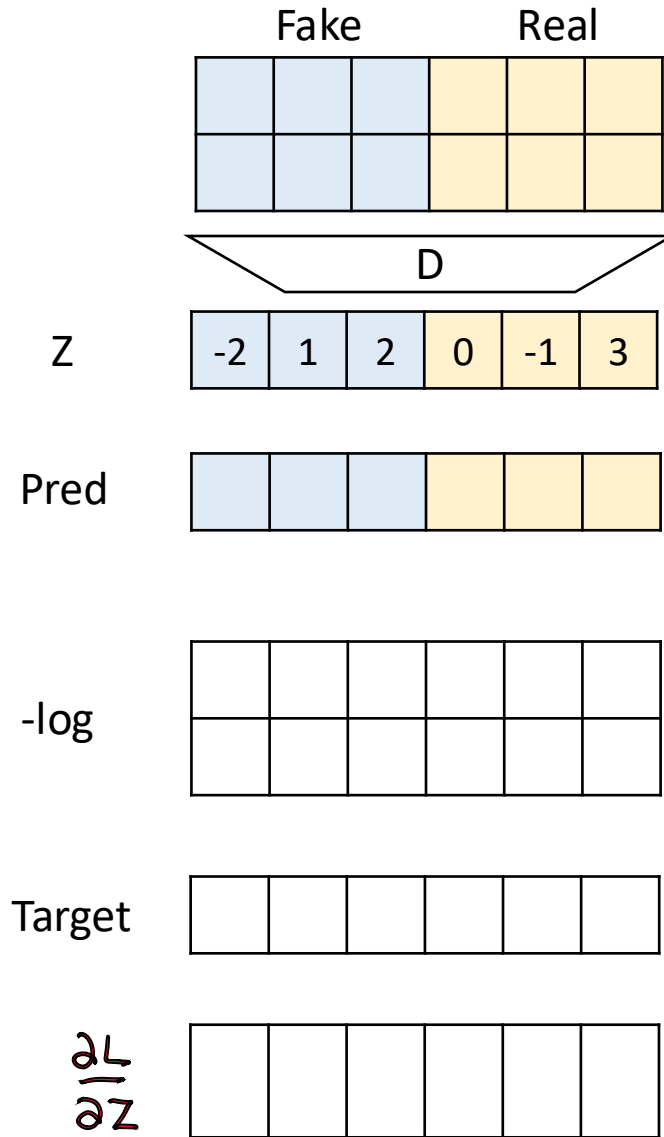
Generator Loss Gradients



$\sigma(\leq -3) \approx 0$
 $\sigma(-2) \approx 0.1$
 $\sigma(-1) \approx 0.3$
 $\sigma(0) \approx 0.5$
 $\sigma(1) \approx 0.7$
 $\sigma(2) \approx 0.9$
 $\sigma(\geq 3) \approx 1$

$-\log(0) \approx 32$
 $-\log(0.1) \approx 3$
 $-\log(0.3) \approx 2$
 $-\log(0.5) \approx 1$
 $-\log(0.7) \approx 0.5$
 $-\log(0.9) \approx 0.2$
 $-\log(1) = 0$

○ Discriminator Loss Gradients



a

🔑

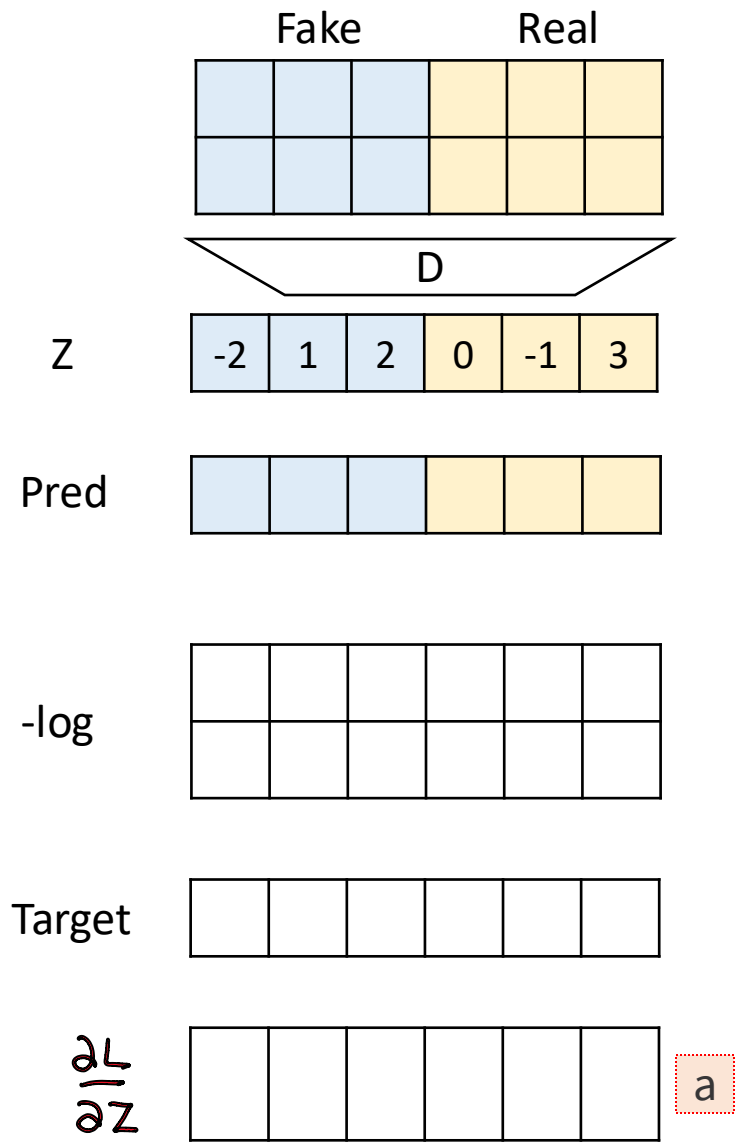
|sum(a)| = .5

$\sigma(\leq -3) \approx 0$
 $\sigma(-2) \approx 0.1$
 $\sigma(-1) \approx 0.3$
 $\sigma(0) \approx 0.5$
 $\sigma(1) \approx 0.7$
 $\sigma(2) \approx 0.9$
 $\sigma(\geq 3) \approx 1$

$-\log(0) \approx 32$
 $-\log(0.1) \approx 3$
 $-\log(0.3) \approx 2$
 $-\log(0.5) \approx 1$
 $-\log(0.7) \approx 0.5$
 $-\log(0.9) \approx 0.2$
 $-\log(1) = 0$




Generator Loss Gradients



$\sigma(\leq -3) \approx 0$
$\sigma(-2) \approx 0.1$
$\sigma(-1) \approx 0.3$
$\sigma(0) \approx 0.5$
$\sigma(1) \approx 0.7$
$\sigma(2) \approx 0.9$
$\sigma(\geq 3) \approx 1$

$-\log(0) \approx 32$
$-\log(0.1) \approx 3$
$-\log(0.3) \approx 2$
$-\log(0.5) \approx 1$
$-\log(0.7) \approx 0.5$
$-\log(0.9) \approx 0.2$
$-\log(1) = 0$

 $|\text{sum}(a)| = 1.3$

Adversarial Training

CSCI 5722 Computer Vision



University of Colorado
Boulder

Train Discriminator

			1	1	-1			
			1	-1	1			
			1	1	1			
			Fake			Real		
1	0	-2	-1	-1	-3	0	0	-1
0	2	1	3	-1	3	3	-1	1
			1	1	1	1	1	1
1	0	1						
0	1	1						
			1	1	1	1	1	1
1	1	-3						
Pred			.3	.2	.4	.5	.3	.7
Target								
Loss Gradients								

Train Generator

The diagram illustrates the forward pass of a Generative Adversarial Network (GAN). It shows the generation of Fake and Real images from latent vectors, the calculation of predicted probabilities, and the target probabilities for the discriminator.

Latent Vectors: A 2x3 grid of latent vectors is shown at the top. The first row contains 1, 1, and -1. The second row contains 1, -1, and 1.

Fake Images: The latent vectors are fed into the Fake generator. The output is a 2x3 grid of Fake images. The first row contains -1, -1, and -3. The second row contains 3, -1, and 3.

Real Images: The latent vectors are also fed into the Real generator. The output is a 2x3 grid of Real images. The first row contains 0, 0, and -1. The second row contains 3, -1, and 1.

Prediction: The Fake and Real images are fed into the discriminator. The output is a 2x3 grid of predicted probabilities. The first row contains .3, .2, and .4. The second row contains .5, .3, and .7.

Target: The target probabilities for the discriminator are shown in a 2x3 grid. The first row contains 1, 1, and 1. The second row contains 0, 0, and 0.

Loss Gradients: The loss gradients are shown in a 2x3 grid. The first row contains 1, 1, and 1. The second row contains 1, 1, and 1.

The figure displays a 3x3 grid of 2D plots showing the evolution of a probability distribution over time. Each plot has a horizontal axis labeled 'x' and a vertical axis labeled 'y'. The plots show a distribution of points (black dots) that changes shape and position from one time step to the next. The top row shows the initial state with a distribution concentrated at x=0. The middle row shows the distribution spreading and moving. The bottom row shows the distribution becoming more complex and spread out. The plots are arranged in a 3x3 grid, with the top row having 3 plots, the middle row having 3 plots, and the bottom row having 3 plots. The plots are labeled with time steps: 0, 1, 2, 3, 4, 5, 6, 7, 8, 9, 10, 11, 12, 13, 14, 15, 16, 17, 18, 19, 20, 21, 22, 23, 24, 25, 26, 27, 28, 29, 30, 31, 32, 33, 34, 35, 36, 37, 38, 39, 40, 41, 42, 43, 44, 45, 46, 47, 48, 49, 50, 51, 52, 53, 54, 55, 56, 57, 58, 59, 60, 61, 62, 63, 64, 65, 66, 67, 68, 69, 70, 71, 72, 73, 74, 75, 76, 77, 78, 79, 80, 81, 82, 83, 84, 85, 86, 87, 88, 89, 90, 91, 92, 93, 94, 95, 96, 97, 98, 99.

The diagram illustrates the training process of a Generative Adversarial Network (GAN) through three iterations. It shows the Generator (Fake) and Discriminator (Real) networks, their inputs, outputs (Predictions and Targets), and the resulting Loss Gradients.

Iteration 1:

- Generator (Fake):** Takes inputs 1, 1, 1 and produces outputs 1, -1, 1.
- Discriminator (Real):** Takes inputs 1, 0, -2 and produces outputs 0, 2, 1.
- Predictions:** .3, .2, .4
- Targets:** (Empty boxes)
- Loss Gradients:** (Empty boxes)

Iteration 2:

- Generator (Fake):** Takes inputs 1, 1, 1 and produces outputs -1, -1, -3.
- Discriminator (Real):** Takes inputs 0, 0, -1 and produces outputs 3, -1, 1.
- Predictions:** (Empty boxes)
- Targets:** (Empty boxes)
- Loss Gradients:** (Empty boxes)

Iteration 3:

- Generator (Fake):** Takes inputs 1, 1, 1 and produces outputs 3, -1, 3.
- Discriminator (Real):** Takes inputs 3, -1, 1 and produces outputs 0, 1, 1.
- Predictions:** (Empty boxes)
- Targets:** (Empty boxes)
- Loss Gradients:** (Empty boxes)

1	0	-2
0	2	1

1.1	0	1
0	1	1

1	1	-3
---	---	----

1	1	-1
1	-1	1

1 1 1

Fake

-1	-1	-3
3	-1	3

1 1 1

--	--	--

.2	.1	.3
----	----	----

--	--	--

--	--	--

--	--	--

1	1	-1
1	-1	1

1	1	1
---	---	---

0	0	-1
3	-1	1

Fake

1	0	-1
0	2	0

0	0	-2
2	-2	2

1	1	1
---	---	---

Real

1	0	1
0	1	1

1	1	1
---	---	---

1	1	1
---	---	---

1	1	-3
---	---	----

1	1	1
---	---	---

1	1	1
---	---	---

.4	.2	.5
----	----	----

.4	.2	.5
----	----	----

.6	.4	.8
----	----	----